

AI Agents in Banking and Revenue Management

A Detailed Scientific Analysis

By V S S Anirudh Sharma

Executive Summary

This whitepaper provides a structured examination of **AI agents** and their role in banking and revenue management. It clarifies fundamental definitions, draws critical distinctions from traditional automation, and evaluates both opportunities and risks in deploying agents in high-stakes financial environments.

- **Conceptual Foundations:** AI agents differ from simple automation by combining perception, reasoning, learning, and adaptive decision-making. They exhibit autonomy, reactivity, proactivity, and social ability, enabling them to operate in dynamic, uncertain environments.
- **Applications in Practice:** Banks already deploy agents in **fraud detection**, **customer service**, **credit risk assessment**, and **billing anomaly detection** within platforms such as Oracle Revenue Management and Billing (ORMB). These examples demonstrate measurable value when coupled with oversight.
- **Constraints and Risks:** Widespread adoption faces hurdles including integration with legacy systems, data quality challenges, regulatory restrictions on autonomy, and accountability gaps. Studies show that **95% of AI projects fail to deliver enterprise value**, underscoring the importance of realistic expectations.
- **Strategic Path Forward:** The paper recommends a **semi-autonomous approach**—embedding human oversight, designing configurable frameworks instead of generic agents, and pursuing phased rollouts with transparent governance and explainability.

Ultimately, AI agents should not aim to replace human judgment but to **augment human capabilities** through accountable, transparent, and context-aware autonomy aligned with the strict demands of financial services.

1. Formal Definitions and Theoretical Framework

1.1 Agent Definition

An **agent** is formally defined as any computational entity that perceives its environment through sensors and acts upon that environment through actuators to achieve specific goals. The canonical formulation by Russell and Norvig establishes four essential properties:^{[2][1]}

1. **Autonomy**: The capacity to operate without constant human intervention, making independent decisions based on internal reasoning mechanisms.
2. **Reactivity**: The ability to perceive environmental changes and respond appropriately in a timely manner.
3. **Proactivity**: Goal-directed behavior that involves taking initiative rather than merely responding to stimuli.
4. **Social Ability**: The capability to interact with other agents, humans, or systems through communication protocols.^{[4][5]}

1.2 Artificial Intelligence (AI)

AI encompasses computational systems that exhibit intelligent behavior, including reasoning, learning, perception, and decision-making capabilities that approximate or exceed human cognitive abilities in specific domains. AI systems employ algorithms, statistical models, and knowledge representations to process information and generate responses that demonstrate intelligence.^[1]

1.3 AI Agent Definition

An **AI agent** represents the convergence of agent theory and artificial intelligence, embodying autonomous systems that utilize AI techniques (machine learning, natural language processing, reasoning algorithms) to perceive, decide, and act in complex environments. AI agents demonstrate rational behavior by selecting actions that maximize expected performance measures given available information and computational constraints.^{[6][7][1]}

2. What Is NOT an Agent: Critical Distinctions

2.1 Computational Systems vs. Agents

Many systems marketed as “agents” today are in fact advanced automation. The distinction matters: mislabeling inflates expectations and undermines genuine innovation.

- **Calculators:** Execute deterministic computations without perception, autonomy, or goals [8][1].
- **Automated Workflows:** Follow predefined sequences without situational awareness or adaptation [8][1].
- **ML Optimizations:** Tune parameters for specific outcomes but lack autonomous reasoning or environmental interaction [1][8].

2.2 The Autonomy Threshold

True agents must demonstrate **decision-making under uncertainty**. This means evaluating multiple options, reasoning about trade-offs, and selecting actions that optimize goals — rather than executing fixed scripts [9][8].

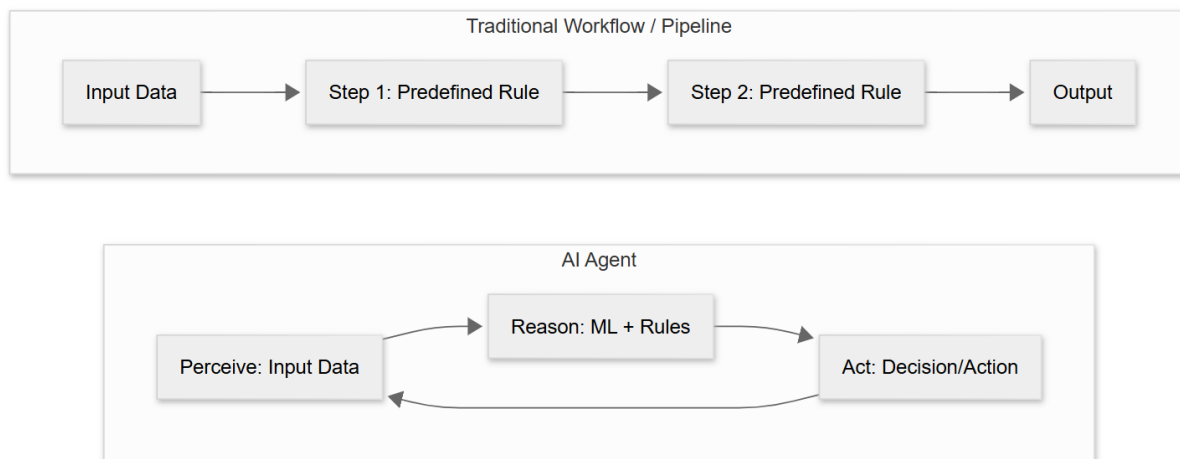


Figure 1: Agent vs Workflow

3. Scope of Agents in Banking: Current Reality and Constraints

3.1 Viable AI Agent Applications

Banking environments accommodate AI agents in specific contexts where accountability frameworks can be maintained —:^{[10][11]}

- **Fraud Detection and Cybersecurity:** AI agents monitor transaction patterns, learn from evolving fraud techniques, and autonomously flag suspicious activities while maintaining audit trails, these systems demonstrate true agency through adaptive learning and proactive threat identification.^{[12][13]}
- **Customer Service and Support:** Conversational agents like Bank of America's Erica serve millions of clients with personalized financial guidance, demonstrating autonomous reasoning and goal-directed assistance.^{[10][12]}
- **Credit Risk Assessment:** AI agents analyze alternative data sources, employment patterns, and behavioral indicators to make lending decisions, exhibiting rational choice under uncertainty.^{[14][15]}
- **Compliance and KYC Monitoring:** Increasingly, agents are applied to Anti-Money Laundering (AML) and Know-Your-Customer (KYC) checks, scanning for anomalies across identity, transaction, and network data in real time.

3.2 Regulatory and Accountability Constraints

Banking's regulatory environment imposes strict limitations on agent autonomy,^{[16][17]}

- **Explainability Requirements:** Financial decisions must be auditable and explainable, constraining the use of "black box" AI systems, The EU AI Act classifies agentic finance tools as "high risk," requiring AI explainability, human controls, and third-party audits.^{[18][19][16]}
- **Liability Frameworks:** Current legal structures struggle to assign responsibility for autonomous AI decisions, necessitating human-in-the-loop architectures for critical financial transactions.^{[17][16]}

- **Data Privacy and Security:** Agents handling sensitive financial data must comply with GDPR, CCPA, and banking-specific regulations, limiting their autonomous access to customer information.^{[20][16]}

3.3 Semi-Autonomous Agent Model

Banking predominantly employs **semi-autonomous agents** that operate within defined boundaries while escalating complex decisions to human operators, This approach balances efficiency gains with regulatory compliance and risk management requirements.^{[5][11]}

4. Scope of Agents in Revenue Management and Billing

4.1 True Agent Applications

Revenue management and billing domains support genuine AI agents through:

- **Dynamic Pricing Optimization:** Agents analyze market conditions, customer behavior, and competitive factors to autonomously adjust pricing strategies while pursuing revenue maximization goals.^{[15][21]}
- **Billing Anomaly Detection:** Enterprise platforms for revenue management and billing, such as ORMB, incorporate AI/ML capabilities for autonomous anomaly detection, where agents learn from historical billing patterns and proactively identify irregularities.^{[22][23]}
- **Revenue Leakage Prevention:** Agents continuously monitor billing processes to identify and address revenue losses through pattern recognition and adaptive learning.^{[24][25]}

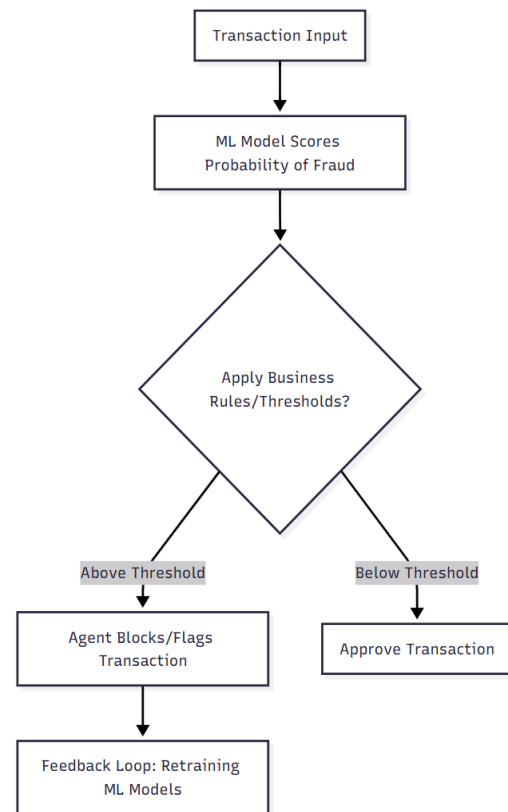


Fig 2: Hybrid ML + Rules Architecture

4.2 Distinguishing Agents from Automation

Critical differentiation factors include:

- **Adaptive Learning:** True agents learn from experience and modify behavior, unlike static workflows.
- **Contextual Decision-Making:** Agents consider multiple variables and trade-offs rather than following predetermined rules.
- **Goal Optimization:** Agents pursue complex objectives (revenue maximization, customer satisfaction, compliance) rather than executing fixed tasks.^{[8][11]}

5. Accountability and High-Stakes Challenges

5.1 The Accountability Paradox

Banking and financial services face a fundamental tension between agent autonomy and accountability requirements.^{[16][20]}

- **Attribution of Responsibility:** When autonomous agents make consequential financial decisions, determining liability becomes complex, particularly when multiple agents interact or when decisions result from emergent behaviors.^{[19][16]}
- **Regulatory Compliance:** Financial institutions must ensure that agent decisions comply with evolving regulations while maintaining transparency and auditability.^{[17][18]}
- **Trust and Transparency:** Customer trust requires explainable decisions, creating tension with sophisticated AI models that may lack interpretability.^{[18][16]}

5.2 Risk Mitigation Strategies

Financial institutions employ several approaches to address accountability challenges:

- **Human-in-the-Loop Architectures:** Critical decisions require human oversight and approval, limiting agent autonomy for high-stakes transactions.^{[11][18]}
- **Audit Trail Requirements:** All agent actions must be logged, traceable, and explainable for regulatory compliance.^{[17][18]}
- **Graduated Autonomy:** Agents operate with varying levels of independence based on decision impact and risk assessment.^{[11][19]}

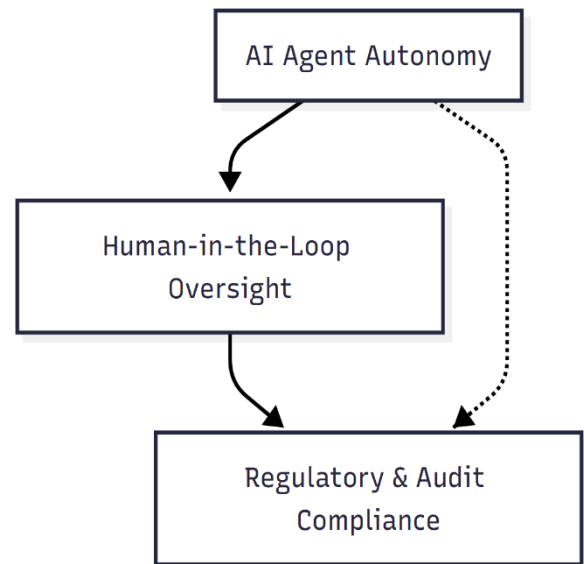


Fig 3: Accountability Layers

6. Feasibility of Generic AI Agent Products for ORMB

6.1 Technical Feasibility

Oracle Revenue Management and Billing (ORMB) demonstrates technical capability for AI agent integration through its existing AI/ML features for billing anomaly detection. The platform's architecture supports:^{[23][22]}

- **Data Integration:** ORMB can aggregate multi-source data necessary for agent perception and decision-making.^{[26][27]}
- **Rule Engine Integration:** Existing business rules can be augmented with AI agent reasoning capabilities.^{[28][24]}
- **API Connectivity:** Modern ORMB implementations support API-based integration enabling agent interaction with external systems.^{[29][30]}

6.2 Customization Requirements

Creating generic AI agent products for ORMB faces significant customization challenges:

- **Heterogeneous Banking Environments:** Banks differ substantially in data formats, business rules, regulatory requirements, and legacy system integrations, necessitating extensive customization.^{[26][16]}
- **Regulatory Variation:** Different jurisdictions impose varying compliance requirements, making truly generic agents impractical without substantial localization.^{[20][16]}
- **Risk and Control Frameworks:** Each institution maintains unique risk tolerance and control mechanisms that must be reflected in agent behavior.^{[19][17]}

6.3 Market Reality Assessment

Industry evidence suggests that successful AI agent deployment in enterprise billing platforms requires:

- **Configurable Agent Frameworks:** Rather than generic agents, successful products provide configurable platforms that can be tailored to specific banking environments.^{[30][31]}
- **Professional Services Integration:** Implementation typically requires substantial consulting and customization services, as evidenced by Deloitte's ORMB accelerator offerings.^[28]
- **Gradual Deployment Models:** Successful implementations follow phased approaches, starting with low-risk applications and gradually expanding agent autonomy.^{[11][19]}

7. Critical Analysis and Limitations

7.1 The Automation-Agent Spectrum

Many systems currently labeled as "AI agents" in banking and billing represent **advanced automation with ML components** rather than true agents exhibiting autonomy, goal-directed behavior, and adaptive learning, this misclassification creates unrealistic expectations and undermines genuine agent development.^{[32][16]}




Fig 4: Agent Spectrum in Banking

7.2 Scalability Challenges

Research indicates that **95% of AI projects fail to deliver enterprise value**, with primary factors including.^{[33][16]}

- **Integration Complexity:** Enterprise systems require extensive customization that undermines generic product viability.
- **Data Quality Issues:** Agents require high-quality, consistent data that many organizations lack.
- **Change Management:** Successful agent deployment requires organizational transformation beyond technology implementation.

7.3 Regulatory Evolution

The regulatory landscape continues evolving, with emerging frameworks like the EU AI Act creating additional compliance requirements for autonomous systems in finance. This regulatory uncertainty complicates long-term agent development strategies.^[19]

8. Conclusions and Recommendations

- **Banking Viability:** AI agents are useful in semi-autonomous roles with human oversight (fraud detection, customer service, anomaly detection).
- **Revenue Management Potential:** Billing and revenue management are stronger candidates for true agents, with lower regulatory burden and clearer KPIs.
- **Generic Agents:** Technically possible, but commercially impractical due to customization, regulatory, and risk-control differences across institutions.

8.4 Strategic Recommendations

- **For Technology Vendors:** Focus on configurable agent frameworks rather than generic agents.
- **For Financial Institutions:** Start with low-risk use cases, embed governance, and expand autonomy gradually.
- **For Regulators:** Define clear accountability frameworks that encourage innovation while protecting consumers and financial stability.

The future of AI agents in banking and revenue management lies not in replacing human judgment but in augmenting human capabilities through carefully designed, accountable, and transparent autonomous systems that respect the unique requirements of financial services domains.

Disclaimer: The views and opinions expressed in this whitepaper are solely those of the author and do not necessarily reflect the official policy, position, or views of Oracle or any of its affiliates. All information about Oracle products (such as Oracle Revenue Management and Billing) is derived from publicly available documentation and sources. This document is for educational and informational purposes only and should not be considered official guidance or endorsement by Oracle.

References

- [1] S. Russell and P. Norvig, *Artificial Intelligence: A Modern Approach*, 3rd ed. Pearson, 2003.
- [2] GeeksforGeeks, “Autonomous agents and their everyday impact,” 2023. [Online]. Available: <https://www.geeksforgeeks.org/artificial-intelligence/autonomous-agents-and-their-everyday-impact/>. [Accessed: Aug. 31, 2025].
- [3] S. Russell and P. Norvig, “PEAS framework,” 2003. [Online]. Available: <https://saiwa.ai/blog/peas-in-ai/>. [Accessed: Aug. 31, 2025].
- [4] Simplilearn, “Exploring intelligent agents and its types, functions & composition,” 2023. [Online]. Available: <https://www.simplilearn.com/peas-in-ai-article>. [Accessed: Aug. 31, 2025].
- [5] IBM, “What are AI agents?,” 2025. [Online]. Available: <https://www.ibm.com/topics/ai-agents>. [Accessed: Aug. 31, 2025].
- [6] S. Russell and P. Norvig, “Rational agents and performance measures,” 2003. [Online]. Available: <https://aima.cs.berkeley.edu/>. [Accessed: Aug. 31, 2025].
- [7] GeeksforGeeks, “Types of AI agents in artificial intelligence,” 2023. [Online]. Available: <https://www.geeksforgeeks.org/artificial-intelligence/agents-artificial-intelligence/>. [Accessed: Aug. 31, 2025].
- [8] GeeksforGeeks, “Understanding PEAS in artificial intelligence,” 2023. [Online]. Available: <https://www.geeksforgeeks.org/artificial-intelligence/understanding-peas-in-artificial-intelligence/>. [Accessed: Aug. 31, 2025].
- [9] Accelirate, “AI agents in banking: The future of smarter, faster, safer finance,” 2025. [Online]. Available: <https://www.accelirate.com/insights/ai-agents-in-banking/>. [Accessed: Aug. 31, 2025].
- [10] Everest Group, “Banking on autonomous agents: Embracing agentic AI in financial services,” 2021. [Online]. Available: <https://www.everestgrp.com/2021-01-banking-on-autonomous-agents.html>. [Accessed: Aug. 31, 2025].

- [11] Akira AI, “Revolutionizing banking operations with AI agents,” 2025. [Online]. Available: <https://www.akira.ai/blog/revolutionizing-banking-operations-with-ai-agents>. [Accessed: Aug. 31, 2025].
- [12] Tredence, “7 best AI agents for financial services,” 2025. [Online]. Available: <https://tredence.com/blog/7-best-ai-agents-for-financial-services/>. [Accessed: Aug. 31, 2025].
- [13] Multimodal, “AI agents in finance: 17 real-world use cases,” 2025. [Online]. Available: <https://multimodal.ai/blog/ai-agents-in-finance-17-real-world-use-cases/>. [Accessed: Aug. 31, 2025].
- [14] Aisera, “16 real-world AI agent examples in 2025,” 2025. [Online]. Available: <https://www.aisera.com/blog/16-real-world-ai-agent-examples-in-2025/>. [Accessed: Aug. 31, 2025].
- [15] Winvesta, “Ethical and regulatory challenges of agentic AI in financial services,” 2025. [Online]. Available: <https://winvesta.in/blog/ethical-and-regulatory-challenges-agentic-ai/>. [Accessed: Aug. 31, 2025].
- [16] SouthState Correspondent, “Why banks must validate the identity of AI agents,” 2025. [Online]. Available: <https://www.southstatecorrespondent.com/why-banks-must-validate-ai-agents/>. [Accessed: Aug. 31, 2025].
- [17] SalesCloser, “How explainable AI is revolutionizing financial services,” 2025. [Online]. Available: <https://www.salescloser.ai/explainable-ai-financial-services/>. [Accessed: Aug. 31, 2025].
- [18] Grant Thornton, “Banks see benefits of AI in regulatory compliance,” 2025. [Online]. Available: <https://www.grantthornton.com/insights/articles/banks-benefits-ai-compliance>. [Accessed: Aug. 31, 2025].
- [19] Infosys, “What agentic AI means for financial services,” 2025. [Online]. Available: <https://www.infosys.com/industries/financial-services/ai-agentic-ai.html>. [Accessed: Aug. 31, 2025].

- [20] Deloitte, “Raise the revenue line with Deloitte digital revenue platform,” 2025. [Online]. Available: <https://www2.deloitte.com/us/en/pages/technology/articles/digital-revenue-platform.html>. [Accessed: Aug. 31, 2025].
- [21] Oracle Docs, “Billing anomaly detection with artificial intelligence (AI) and ML,” 2025. [Online]. Available: https://docs.oracle.com/en/cloud/saas/billing-cloud/csbi_matching_AI.html. [Accessed: Aug. 31, 2025].
- [22] Oracle Docs, “Oracle revenue management and billing,” 2025. [Online]. Available: <https://docs.oracle.com/en/cloud/saas/revenue-management-billing/index.html>. [Accessed: Aug. 31, 2025].
- [23] Fusion Practices, “ORMB integration articles,” 2025. [Online]. Available: <https://fusionpractices.com/articles/oracle-revenue-management-billing-introduction/>. [Accessed: Aug. 31, 2025].
- [24] SunTec, “Enterprise billing & pricing management solutions,” 2025. [Online]. Available: <https://www.suntecgroup.com/billing-solution/>. [Accessed: Aug. 31, 2025].
- [25] Oracle Docs, “Oracle revenue management and billing: Data integration,” 2025. [Online]. Available: <https://docs.oracle.com/en/cloud/saas/revenue-management-billing/data-integration.html>. [Accessed: Aug. 31, 2025].
- [26] Oracle Docs, “Oracle revenue management billing architecture,” 2025. [Online]. Available: <https://docs.oracle.com/en/cloud/saas/revenue-management-billing/architecture-overview.html>. [Accessed: Aug. 31, 2025].
- [27] Deloitte, “Deloitte digital revenue platform accelerator,” 2025. [Online]. Available: <https://www2.deloitte.com/us/en/pages/deloitte-digital/articles/revenue-platform-accelerator.html>. [Accessed: Aug. 31, 2025].
- [28] GetMonetizely, “Integration platforms for AI agent billing,” 2025. [Online]. Available: <https://www.getmonetizely.com/platform/integration>. [Accessed: Aug. 31, 2025].

- [29] J.P. Morgan, “BillingPlatform partner integration,” 2025. [Online]. Available: <https://www.jpmorgan.com/solutions/cib/billing-platform>. [Accessed: Aug. 31, 2025].
- [30] Oracle Docs, “Cloud service API connectivity,” 2025. [Online]. Available: <https://docs.oracle.com/en/cloud/saas/api-connectivity/>. [Accessed: Aug. 31, 2025].
- [31] NayaOne, “Decoding AI: Automation, workflows and agents,” 2025. [Online]. Available: <https://nayaone.com/blog/decoding-ai-automation-workflows-agents/>. [Accessed: Aug. 31, 2025].
- [32] Fortune, “Why 95% of AI projects fail to deliver enterprise value (MIT report),” 2025. [Online]. Available: <https://fortune.com/2025/08/18/mit-report-95-percent-generative-ai-pilots-at-companies-failing-cfo/>. [Accessed: Aug. 31, 2025].
- [33] H. Chip, “Agents,” *Huyenchip Blog*, Jan. 2025. [Online]. Available: <https://huyenchip.com/2025/01/07/agents.html>. [Accessed: Aug. 31, 2025].